


0	0	0	0	0	0
0	1	0	0	1	0
0	0	0	0	0	0
1	0	0	0	0	1
0	1	1	1	1	0
0	0	0	0	0	0

# Can we use traditional feed forward nets?



How many weights and biases do we have (if we have one hidden layer with 36 neurons)?

- Hidden layer:
- Output layer:
- Total:
- **What if the image was of a different size?**
- **Let's take a  $224*224*3$  image**
- Input values:  $224*224*3 = 150\_528$
- Hidden layer:
- Output layer:
- Total:
- This will lead to two problems:

-

# Convolutional Neural Networks


- Don't flatten the image
- Apply convolutional kernels/filters to the image

Input Image

Kernel/Filter

Feature Maps

Step size is defined by the „stride“



# Convolutional Neural Networks

- Don't flatten the image
- Apply convolutional kernels/filters (randomly initialized in the beginning) to the image -> Feature Map
- Add a bias (also randomly initialized in the beginning) to every value of the feature map
- Max Pooling: Slide a window over your image

Input Image

Kernel/Filter

Feature Maps

Add bias

Max Pooling

-1	-1	-1
-1	1	-1
-1	-1	-1

1	-1	-1	1
-2	-1	-1	-2
-3	-3	-3	-3
-1	-1	-1	-1

b = 0

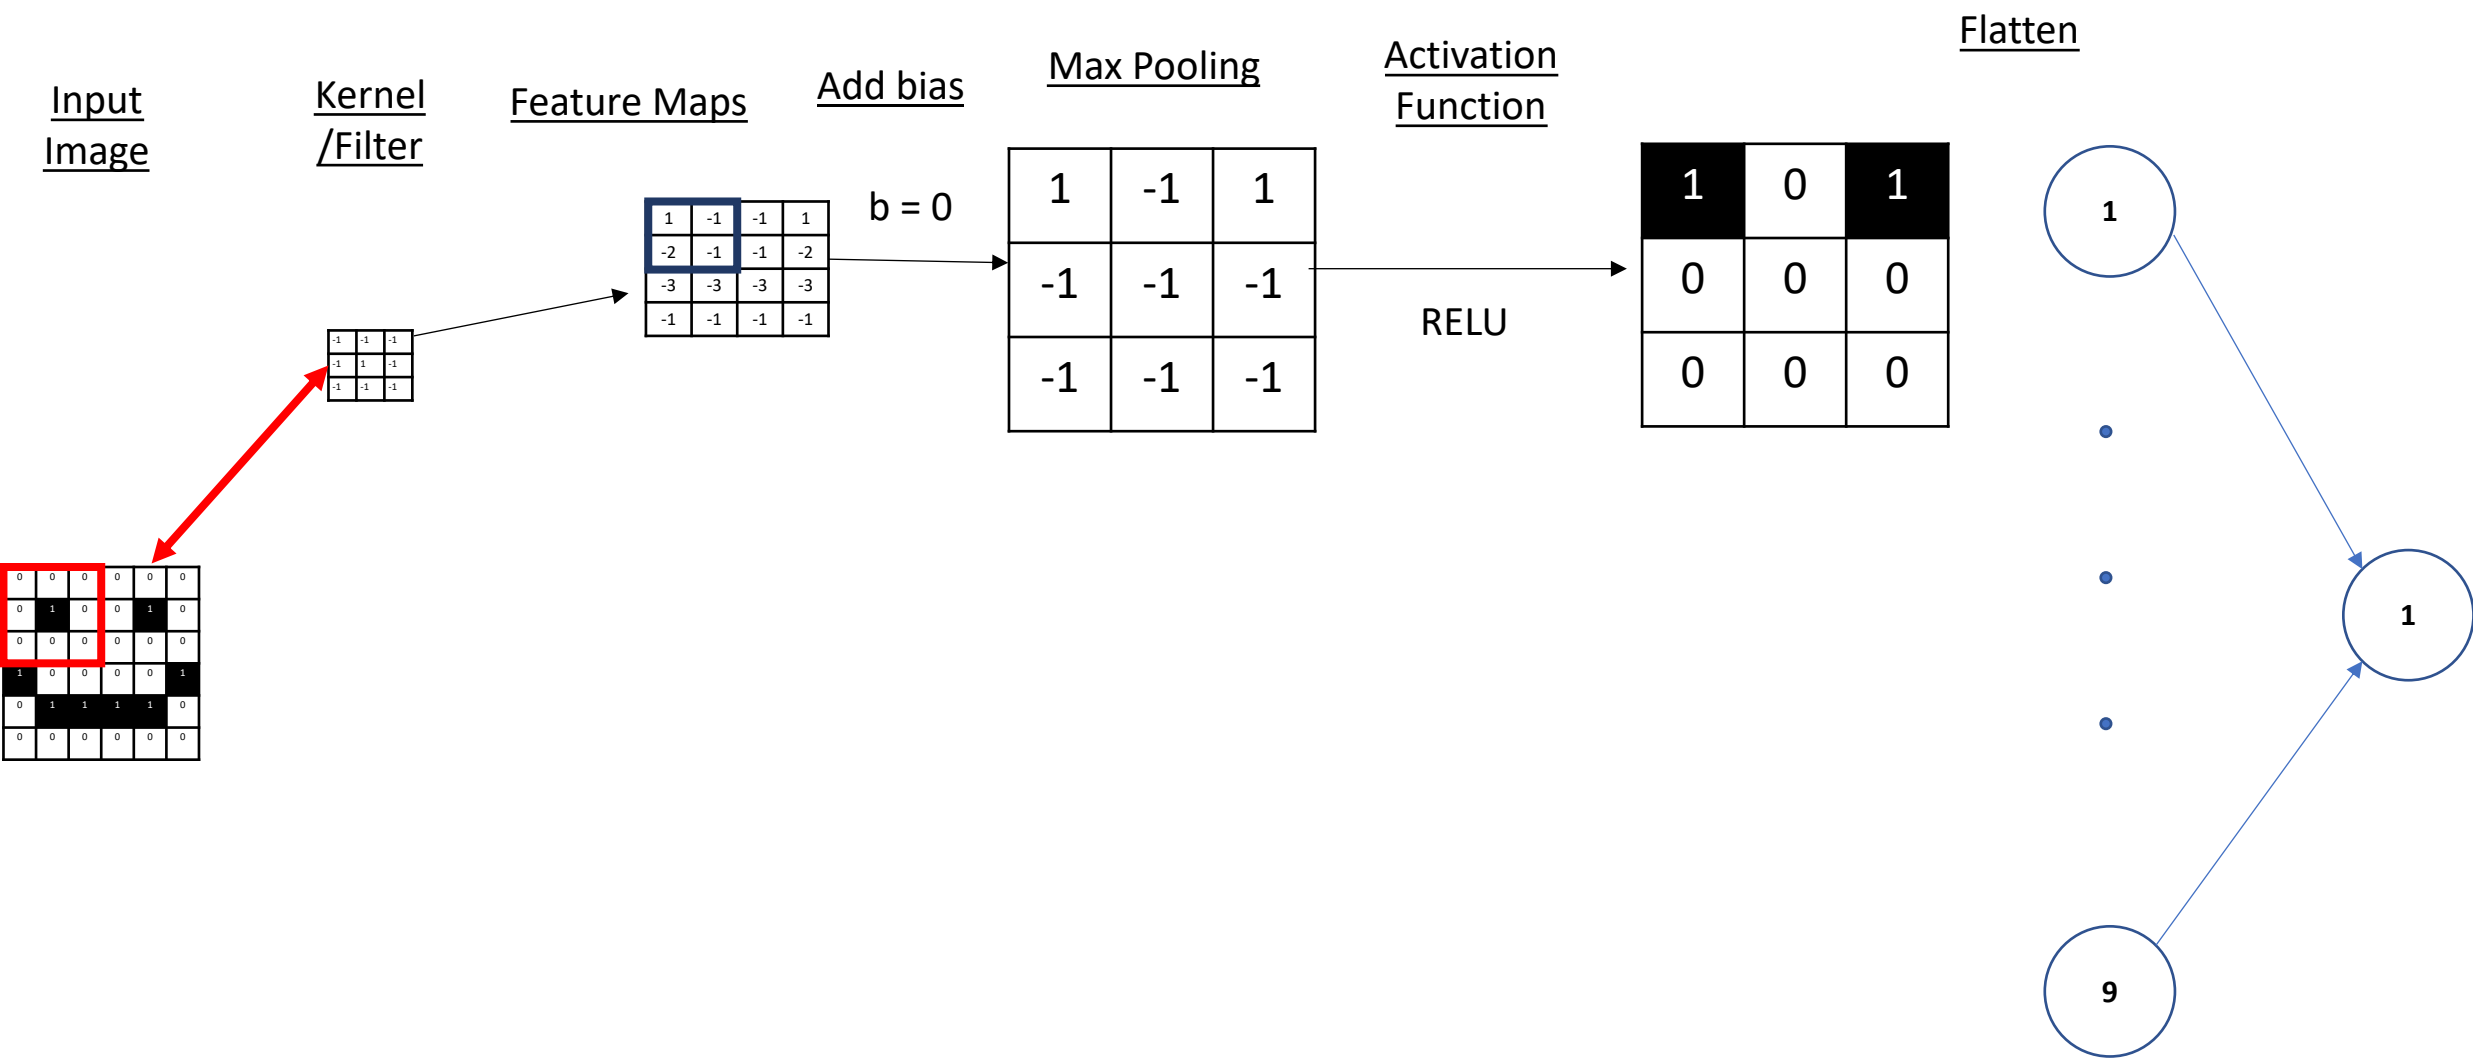
1	-1	1
-1	-1	-1
-1	-1	-1

0	0	0	0	0	0
0	1	0	0	1	0
0	0	0	0	0	0
1	0	0	0	0	1
0	1	1	1	1	0
0	0	0	0	0	0

# Convolutional Neural Networks

- Don't flatten the image
- Apply convolutional kernels/filters (randomly initialized in the beginning) to the image -> Feature Map
- Add a bias (also randomly initialized in the beginning) to every value of the feature map
- Max Pooling: Slide a window over your image → End up with a reduced feature map
- Activation Function: In our example we will use ReLu





Input Image

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
1	0	0	0	0	0	1
0	1	1	1	1	1	0
0	0	0	0	0	0	0

Convolutional  
Kernel

-1	-1	-1
-1	1	-1
-1	-1	-1

Add bias

$b = 0$

Feature Map

1	-1	-1	1
-2	-1	-1	-2
-3	-3	-3	-3
-1	-1	-1	-1

Max Pooling

1	-1	1
-1	-1	-1
-1	-1	-1

ReLu/Activation

1	0	1
0	0	0
0	0	0

$b = -2$

1	-1	1
1	1	1
-1	-1	-1

-1	-1	-1	-1
-4	-1	-1	-4
-3	-5	-5	-3
1	1	1	1

-1	-1	-1
-1	-1	-1
1	1	1

0	0	0
0	0	0
1	1	1

...

ReLu/Activation

1	0	1
0	0	0
0	0	0

0	0	0
0	0	0
1	1	1

Flatten

1
0
1
0
0
0
0
0
0
0
0
0
0
1
1
1

Dense

...

$f(x)$

# Convolutional Layers

- They also have weights and biases. The number just does not increase with the image size.
- Also, convolutional layers exploit the fact that pixels are meaningful in their close environment, not by themselves.
- Do Convolutional Layers suffer from problems?
- Yes:
  - Filter reduces the size of the feature space → loss of information at the edges
  - Solution: padding

# Convolutional Layers - Hyperparameters

- What are hyperparameters of convolutional layers?
  - Padding? – yes or no
  - Nr. of kernels
  - Size of kernels, 3x3, 5x5
  - Stride of kernels, (1x1), (2x2)
  - Size of Max-Pooling layer, 2x2, 3x3
  - Stride of Max-Pooling layer, 2x2
  - Activation Function, ReLu
  - Nr. of convolutional layers
  - Sequence of layer parts

# Convolutional Layers – Max-Pooling

- There are two main ideas behind Max-Pooling:
  - 1) Reduce the number of features
  - 2) Small variations (shifts by a few pixels) get cancelled out

# Convolutional Layers – Padding

- Padding ensures that you don't lose information at the edge of the image

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	1	0	0	1	0	0
0	0	0	0	0	0	0	0
0	1	0	0	0	0	1	0
0	0	1	1	1	1	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0